

Kubernetes

Arun

May 27, 2019



kubernetes

First bit of basics



kubernetes

First bit of basics

- pods
- deployment
- service
- horizontal pod autoscaler (HPA)



kubernetes

Before k8s



kubernetes

Before k8s

- 600+ servers
- External load balancers see 5K+ requests per *second*
- Internal Amplification of 8x to 12x



kubernetes

Kubernetes



kubernetes

- Partial Kubernetes deployment since Oct, 2016
- Full Production deployment since Nov, 2016
- Using Google Container Engine



kubernetes

Kubernetes



kubernetes

- Partial Kubernetes deployment since Oct, 2016
- Full Production deployment since Nov, 2016
- Using Google Container Engine
- 5+ Clusters
- 75+ deployments
- 500+ containers (At Peak)



kubernetes



kubernetes

- Partial Kubernetes deployment since Oct, 2016
- Full Production deployment since Nov, 2016
- Using Google Container Engine
- 5+ Clusters
- 75+ deployments
- 500+ containers (At Peak)
- All new services directly onboard to kubernetes
- Autoscaled on CPU targets



kubernetes

Kubernetes

```
1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "159"
6     kubernetes.io/change-cause: kubectl set image deployment/main-app main-app=xxxxxxxxxxxxx
7     --record=true
8   creationTimestamp: 2017-05-11T08:17:58Z
9   generation: 3757
10  labels:
11    app: main-app
12    track: stable
13  name: main-app
14  namespace: default
15  resourceVersion: "4890666"
16  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/main-app
17  uid: 57d8c2f8-3622-11e7-87c9-420096
18 spec:
19   replicas: 212
20   revisionHistoryLimit: 15
21   selector:
22     matchLabels:
23       app: main-app
24       track: stable
25   strategy:
26     rollingUpdate:
27       maxSurge: 20%
28       maxUnavailable: 5%
29     type: RollingUpdate
30   template:
31     metadata:
32       creationTimestamp: null
33     labels:
34       app: main-app
35       track: stable
36     spec:
37       containers:
38       - args:
39         - start_prod
40         env:
41         - name: DISABLE_CONSUL
42           value: "true"
43         - name: SERVICE_NAME
44           value: app
45         - name: GINICORN_THREADS
46           value: "14"
47         - name: DISABLE_STATSD_EXPORTER
48           value: "true"
49         image: xxxxxxxxxxxx
50         imagePullPolicy: Always
51         livenessProbe:
52           failureThreshold: 3
53           httpGet:
54             path: /healthcheck/
55             port: 8080
56             scheme: HTTP
```

● Rolling updates across containers



kubernetes

Kubernetes

```
1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "159"
6     kubernetes.io/change-cause: kubectl set image deployment/main-app main-app=xxxxxxxxxxxxx
7     --record=true
8   creationTimestamp: 2017-05-11T08:17:58Z
9   generation: 3757
10  labels:
11    app: main-app
12    track: stable
13  name: main-app
14  namespace: default
15  resourceVersion: "4890666"
16  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/main-app
17  uid: 57d8c2f8-3622-11e7-87c9-420096
18 spec:
19   replicas: 212
20   revisionHistoryLimit: 15
21   selector:
22     matchLabels:
23       app: main-app
24       track: stable
25   strategy:
26     rollingUpdate:
27       maxSurge: 20%
28       maxUnavailable: 5%
29     type: RollingUpdate
30   template:
31     metadata:
32       creationTimestamp: null
33     labels:
34       app: main-app
35       track: stable
36   spec:
37     containers:
38       - args:
39         - start_prod
40         env:
41         - name: DISABLE_CONSU-
42           value: "true"
43         - name: SERVICE_NAME
44           value: app
45         - name: GOMAXPROCS_THREADS
46           value: "14"
47         - name: DISABLE_STATSD_EXPORTER
48           value: "true"
49         image: xxxxxxxxxxxx
50         imagePullPolicy: Always
51         livenessProbe:
52           failureThreshold: 3
53           httpGet:
54             path: /healthcheck/
55             port: 8080
56             scheme: HTTP
```

- Rolling updates across containers
- Custom commands to run same container across dev/stage/prod environments



kubernetes

Kubernetes

```
1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "159"
6     kubernetes.io/change-cause: kubectl set image deployment/main-app main-app=xxxxxxxxxxxxx
7     --record=true
8   creationTimestamp: 2017-05-11T08:17:58Z
9   generation: 3757
10  labels:
11    app: main-app
12    track: stable
13  name: main-app
14  namespace: default
15  resourceVersion: "4890666"
16  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/main-app
17  uid: 57d8c2f8-3622-11e7-87c9-420096
18 spec:
19   replicas: 212
20   revisionHistoryLimit: 15
21   selector:
22     matchLabels:
23       app: main-app
24       track: stable
25   strategy:
26     rollingUpdate:
27       maxSurge: 20%
28       maxUnavailable: 5%
29     type: RollingUpdate
30   template:
31     metadata:
32       creationTimestamp: null
33     labels:
34       app: main-app
35       track: stable
36     spec:
37       containers:
38       - args:
39         - start_prod
40         env:
41         - name: DISABLE_CONSU
42           value: "true"
43         - name: SERVICE_NAME
44           value: app
45         - name: QUICORN_THREADS
46           value: "14"
47         - name: DISABLE_STATSD_EXPORTER
48           value: "true"
49       image: xxxxxxxxxxxx
50       imagePullPolicy: Always
51       livenessProbe:
52         failureThreshold: 3
53         httpGet:
54           path: /healthcheck/
55           port: 8080
56           scheme: HTTP
```

- Rolling updates across containers
- Custom commands to run same container across dev/stage/prod environments
- Environment variables to override 'config-service' values



kubernetes

Kubernetes


```
1 apiVersion: extensions/v1beta1
2 kind: Deployment
3 metadata:
4   annotations:
5     deployment.kubernetes.io/revision: "159"
6     kubernetes.io/change-cause: kubectl set image deployment/main-app main-app=xxxxxxxxxxxxx
7     --record=true
8   creationTimestamp: 2017-05-11T08:17:58Z
9   generation: 3757
10  labels:
11    app: main-app
12    track: stable
13  name: main-app
14  namespace: default
15  resourceVersion: "4890666"
16  selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/main-app
17  uid: 57d8c2f8-3622-11e7-87c9-420096
18 spec:
19   replicas: 212
20   revisionHistoryLimit: 15
21   selector:
22     matchLabels:
23       app: main-app
24       track: stable
25   strategy:
26     rollingUpdate:
27       maxSurge: 20%
28       maxUnavailable: 5%
29     type: RollingUpdate
30   template:
31     metadata:
32       creationTimestamp: null
33     labels:
34       app: main-app
35       track: stable
36     spec:
37       containers:
38       - args:
39         - start_prod
40         env:
41         - name: DISABLE_CONSUL
42           value: "true"
43         - name: SERVICE_NAME
44           value: app
45         - name: GINICORN_THREADS
46           value: "14"
47         - name: DISABLE_STATSD_EXPORTER
48           value: "true"
49         image: xxxxxxxxxxxx
50         lifecycle:
51           always:
52             failureThreshold: 3
53             httpGet:
54               path: /healthcheck/
55               port: 8080
56             scheme: HTTP
```

- Rolling updates across containers
- Custom commands to run same container across dev/stage/prod environments
- Environment variables to override 'config-service' values
- Liveness and Readiness checks are must!



kubernetes

Kubernetes Autoscaling



kubernetes

Kubernetes

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
chat-app	30	30	30	30	75d
main-app	204	204	204	204	75d
search-app	36	36	36	36	75d
workers-accounts	9	9	9	9	22d
workers-accounts-auth	1	1	1	1	22d
workers-activities	1	1	1	1	22d
workers-apns	1	1	1	1	22d
workers-apns-feedback	1	1	1	1	22d
workers-beat	1	1	1	1	22d
workers-bulk-index-products	4	4	4	4	22d
workers-dump	4	4	4	4	22d
workers-chat-emails	10	10	10	10	22d
workers-message	3	3	3	3	22d
workers-comments	3	3	3	3	22d
workers-delay-images	5	5	5	5	22d
workers-duplicate-products	6	6	6	6	22d
workers-emailalias	1	1	1	1	22d
workers-fblists	5	5	5	5	22d
workers-fbshare	11	11	11	11	22d
workers-flaggings	2	2	2	2	22d
workers-grouplistings	2	2	2	2	22d
workers-listings	5	5	5	5	22d
workers-market	1	1	1	1	22d
workers-purchases	1	1	1	1	22d
workers-transactions	1	1	1	1	22d
workers-notifications	9	9	9	9	10d
workers-notifications-1	3	3	3	3	22d
workers-password	1	1	1	1	22d
workers-post-category-check	2	2	2	2	22d
workers-post-price-check	2	2	2	2	22d
workers-price-consumer	1	1	1	1	22d
workers-product-history	2	2	2	2	22d
workers-product-updates-stream	4	4	4	4	22d
workers-producttk	7	7	7	7	22d
workers-push-apns	3	3	3	3	22d
workers-push-notification	2	2	2	2	22d
workers-push-gcm	8	8	8	8	22d
workers-register-android	3	3	3	3	22d
workers-register-ios	5	5	5	5	22d
workers-rnconsumer	1	1	1	1	22d
workers-sls	4	4	4	4	22d
workers-scheduler	1	1	1	1	22d
workers-seo	1	1	1	1	22d
workers-services-fulfill	1	1	1	1	22d
workers-services-visibility	2	2	2	2	22d
workers-signals-handler	4	4	4	4	22d
workers-listings	4	4	4	4	22d
workers-social	1	1	1	1	22d
workers-support	6	6	6	6	22d
workers-verification	1	1	1	1	22d
workers-webhook	2	2	2	2	22d
workers-worker	10	10	10	10	22d

- Not using K8S Ingress/Service



kubernetes

Kubernetes

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
chat-app	30	30	30	30	75d
main-app	204	204	204	204	75d
search-app	36	36	36	36	75d
workers-accounts	9	9	9	9	22d
workers-accounts-auth	1	1	1	1	22d
workers-activities	1	1	1	1	22d
workers-apns	1	1	1	1	22d
workers-apns-feedback	1	1	1	1	22d
workers-beat	1	1	1	1	22d
workers-bulk-index-products	4	4	4	4	22d
workers-dump	4	4	4	4	22d
workers-chat-emails	10	10	10	10	22d
workers-message	3	3	3	3	22d
workers-comments	3	3	3	3	22d
workers-delay-images	5	5	5	5	22d
workers-duplicate-products	6	6	6	6	22d
workers-emails-las	1	1	1	1	22d
workers-fblists	5	5	5	5	22d
workers-fbshare	11	11	11	11	22d
workers-flaggings	2	2	2	2	22d
workers-group-listings	2	2	2	2	22d
workers-listings	5	5	5	5	22d
workers-market	1	1	1	1	22d
workers-purchases	1	1	1	1	22d
workers-transactions	1	1	1	1	22d
workers-notifications	9	9	9	9	10d
workers-notifications-1	3	3	3	3	22d
workers-password	1	1	1	1	22d
workers-post-category-check	2	2	2	2	22d
workers-post-price-check	2	2	2	2	22d
workers-price-consumer	1	1	1	1	22d
workers-product-history	2	2	2	2	22d
workers-product-updates-stream	4	4	4	4	22d
workers-product-lik	7	7	7	7	22d
workers-push-apns	3	3	3	3	22d
workers-push-notification	2	2	2	2	22d
workers-push-gcm	8	8	8	8	22d
workers-register-android	3	3	3	3	22d
workers-register-ios	5	5	5	5	22d
workers-rn-consumer	1	1	1	1	22d
workers-sls	4	4	4	4	22d
workers-scheduler	1	1	1	1	22d
workers-seo	1	1	1	1	22d
workers-services-fulfill	1	1	1	1	22d
workers-services-visibility	2	2	2	2	22d
workers-signals-handler	4	4	4	4	22d
workers-listings	4	4	4	4	22d
workers-social	1	1	1	1	22d
workers-support	6	6	6	6	22d
workers-verification	1	1	1	1	22d
workers-wback	2	2	2	2	22d
workers-worker	10	10	10	10	22d

- Not using **K8S Ingress/Service**
- Config-Service as DaemonSet
- Containers get registered on Config-Service (*NodePort*) from health check



kubernetes

Kubernetes

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
chat-app	30	30	30	30	75d
main-app	204	204	204	204	75d
search-app	36	36	36	36	75d
workers-accounts	9	9	9	9	22d
workers-accounts-auth	1	1	1	1	22d
workers-activities	1	1	1	1	22d
workers-apns	1	1	1	1	22d
workers-apns-feedback	1	1	1	1	22d
workers-beat	1	1	1	1	22d
workers-bulk-index-products	4	4	4	4	22d
workers-dump	4	4	4	4	22d
workers-chat-emails	10	10	10	10	22d
workers-message	3	3	3	3	22d
workers-comments	3	3	3	3	22d
workers-delay-images	5	5	5	5	22d
workers-duplicate-products	6	6	6	6	22d
workers-emails	1	1	1	1	22d
workers-fblists	5	5	5	5	22d
workers-fbshare	11	11	11	11	22d
workers-flaggings	2	2	2	2	22d
workers-groupings	2	2	2	2	22d
workers-listings	5	5	5	5	22d
workers-market	1	1	1	1	22d
workers-purchases	1	1	1	1	22d
workers-transactions	1	1	1	1	22d
workers-notifications	9	9	9	9	10d
workers-notifications-1	3	3	3	3	22d
workers-password	1	1	1	1	22d
workers-post-category-check	2	2	2	2	22d
workers-post-price-check	2	2	2	2	22d
workers-price-consumer	1	1	1	1	22d
workers-product-history	2	2	2	2	22d
workers-product-updates-stream	4	4	4	4	22d
workers-productlik	7	7	7	7	22d
workers-push-apns	3	3	3	3	22d
workers-push-notification	2	2	2	2	22d
workers-push-gcm	8	8	8	8	22d
workers-register-android	3	3	3	3	22d
workers-register-ios	5	5	5	5	22d
workers-rn-consumer	1	1	1	1	22d
workers-sls	4	4	4	4	22d
workers-scheduler	1	1	1	1	22d
workers-seo	1	1	1	1	22d
workers-services-fulfill	1	1	1	1	22d
workers-services-visibility	2	2	2	2	22d
workers-signals-handler	4	4	4	4	22d
workers-listings	4	4	4	4	22d
workers-social	1	1	1	1	22d
workers-support	6	6	6	6	22d
workers-verification	1	1	1	1	22d
workers-wback	2	2	2	2	22d
workers-worker	10	10	10	10	22d

- Not using **K8S Ingress/Service**
- Config-Service as DaemonSet
- Containers get registered on Config-Service (*NodePort*) from health check
- No change in existing architecture needed
- Service discovery from Internal/External load balancers



kubernetes

Kubernetes

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
chat-app	Deployment/chat-app	48% / 48%	6	90	30	75d
main-app	Deployment/main-app	50% / 48%	18	350	204	75d
search-app	Deployment/search-app	41% / 43%	6	90	36	74d
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
main-app	204	204	204	204	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
chat-app	30	30	30	30	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
search-app	36	36	36	36	75d	

- 'Config-Service' allows us to have hybrid model
- Services can be deployed in multiple clusters
- Instance groups can coexist with Kubernetes



kubernetes

Kubernetes

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
chat-app	Deployment/chat-app	48% / 48%	6	90	30	75d
main-app	Deployment/main-app	50% / 48%	18	350	204	75d
search-app	Deployment/search-app	41% / 43%	6	90	36	74d
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
main-app	204	204	204	204	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
chat-app	30	30	30	30	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
search-app	36	36	36	36	75d	

- 'Config-Service' allows us to have hybrid model
- Services can be deployed in multiple clusters
- Instance groups can coexist with Kubernetes
- Instances run the same container image



kubernetes

Kubernetes

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
chat-app	Deployment/chat-app	48% / 48%	6	90	30	75d
main-app	Deployment/main-app	50% / 48%	18	350	204	75d
search-app	Deployment/search-app	41% / 43%	6	90	36	74d
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
main-app	204	204	204	204	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
chat-app	30	30	30	30	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
search-app	36	36	36	36	75d	

- 'Config-Service' allows us to have hybrid model
- Services can be deployed in multiple clusters
- Instance groups can coexist with Kubernetes
- Instances run the same container image
- Recovery mechanism / Transitioning



kubernetes

Kubernetes

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
chat-app	Deployment/chat-app	48% / 48%	6	90	30	75d
main-app	Deployment/main-app	50% / 48%	18	350	204	75d
search-app	Deployment/search-app	41% / 43%	6	90	36	74d
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
main-app	204	204	204	204	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
chat-app	30	30	30	30	75d	
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE	
search-app	36	36	36	36	75d	

- 'Config-Service' allows us to have hybrid model
- Services can be deployed in multiple clusters
- Instance groups can coexist with Kubernetes
- Instances run the same container image
- Recovery mechanism / Transitioning
- Easy Migration - Create new K8S cluster and terminate old cluster



kubernetes



- Envoy for communication between micro-services (gRPC)



Kubernetes




- Envoy for communication between micro-services (gRPC)
- Prometheus for monitoring



kubernetes

Kubernetes

Hystrix Stream: Auth

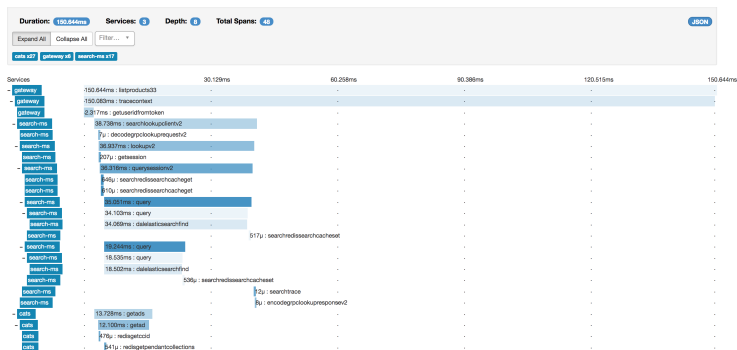


- Envoy for communication between micro-services (gRPC)
- Prometheus for monitoring
- Turbine for hystrix stream aggregation



kubernetes

Kubernetes



- Envoy for communication between micro-services (gRPC)
- Prometheus for monitoring
- Turbine for hystrix stream aggregation
- Zipkin Spans are written to external kafka



kubernetes

Deployment Pipeline

			Build Docker Image	Deploy Image To Canary	Promote Image To Latest	Deploy to Prod	Deploy to Kubernetes	Deploy to Workers
Average stage times: (Average full run time: ~30min 23s)			52s	53s	13s	55s	8min 28s	44s
#1027	Mar 05 15:22	No Changes	57s	47s (paused for 34s)	14s (paused for 7min 0s)	56s	6min 38s	45s
#1026	Mar 05 14:59	No Changes	1min 1s	49s (paused for 2min 47s)	12s (paused for 8min 51s)	1min 1s	6min 30s	41s
#1025	Mar 05 14:05	No Changes	50s	52s (paused for 9min 9s)	14s (paused for 30min 32s)	58s	7min 48s	45s
#1024	Mar 05 12:32	No Changes	45s	49s (paused for 1min 49s)	15s (paused for 11min 48s)	53s	10min 1s	37s
#1023	Mar 02 17:12	No Changes	54s	51s (paused for 22s)	10s (paused for 7min 11s)	52s	9min 8s	44s
#1022	Mar 02 15:40	No Changes	51s	50s (paused for 8s)	15s (paused for 35min 44s)	53s	8min 25s	41s

kubernetes

Thank You

Q&A



kubernetes